

SWE3004 Operating Systems, Fall 2024

Project I. System call

TA)

Osama Khan

Gwanjong Park

Chanu Yu

ShinHyun Park

Project plan

- Total 6 projects
 - 0) Booting xv6 operating system
 - 1) System call
 - 2) CPU scheduling
 - 3) Virtual memory
 - 4) Page replacement
 - 5) File systems

Project 1: Make system calls

- Make system calls in xv6 kernel
 - getnice
 - setnice
 - ps

Project 1: Make system calls

- Goal : make new three system calls(getnice, setnice, ps)
- Synopsis
 - int getnice(int pid);
 - int setnice(int pid, int value);
 - void ps(int pid);
- Description
 - The ***getnice*** function obtains the nice value of a process.
 - The ***setnice*** function sets the nice value of a process.
 - The **default nice value is 20**. Lower nice values cause more favorable scheduling.
 - It will be necessary to implement the nice value before creating the system call.
 - **The range of valid nice value is 0~39**

Project 1: Make system calls

- **Description (cont'd)**
 - In kernel, the **ps** system call prints out process(s)'s information, which includes name, pid, state and priority(nice value) of each process.
 - If the pid is 0, print out all processes' information.
 - Otherwise, print out corresponding process's information.
 - If there is no process corresponding to the pid, print out nothing.

name	pid	state	priority
init	1	SLEEPING	20
sh	2	SLEEPING	20
ps	3	RUNNING	20

- **Return value**
 - **getnice** : Return the nice value of target process on success. Return -1 if there is no process corresponding to the pid.
 - **setnice** : Return 0 on success. Return -1 if there is no process corresponding to the pid or the nice value is invalid.
 - **ps** : No return value.

How to add system call (ex. getpname)

I. Add your syscall to usys.S

```
SYSCALL(getpid)
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(getpname)
```

2. Add syscall number to syscall.h

```
#define SYS_link    19
#define SYS_mkdir   20
#define SYS_close   21
#define SYS_getpname 22
```

3. Add extern and syscall element in syscall.c

```
extern int sys_wait(void);
extern int sys_write(void);
extern int sys_uptime(void);
extern int sys_getpname(void);
```

```
[SYS_link]    sys_link,
[SYS_mkdir]   sys_mkdir,
[SYS_close]   sys_close,
[SYS_getpname] sys_getpname,
};
```

How to add system call (ex. getpname)

4. Add a sys_function to sysproc.c

```
int
sys_getpname(void)
{
    int pid;

    if(argint(0, &pid) < 0)
        return -1;
    return getpname(pid);
}
```

5. Add a function that performs a real action to proc.c

```
int
getpname(int pid)
{
    struct proc *p;

    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->pid == pid){
            cprintf("%s\n", p->name);
            release(&ptable.lock);
            return 0;
        }
    }
    release(&ptable.lock);
    return -1;
}
```

How to add system call (ex. getpname)

6. Add a definition to defs.h and user.h

```
void          userinit(void);
int           wait(void);
void          wakeup(void*);
void          yield(void);
int           getpname(int);
```

```
char* sbrk(int);
int sleep(int);
int uptime(void);
int getpname(int);
```

How to test your system call

```
#include "types.h"
#include "user.h"
#include "stat.h"

int main()
{
    int i;
    for (i=1; i<11; i++) {
        printf(1, "%d: ", i);
        if (getpid(i))
            printf(1, "Wrong pid\n");
    }
    exit();
}
```

mytest.c

“mytest.c” is an example code.
Create and use your own test code.

```
UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _zombie\
    _mytest\
```

Makefile

Test with user program

```
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive fi
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ ls
.
..
README 2 2 2286
cat 2 3 13648
echo 2 4 12656
forktest 2 5 8092
grep 2 6 15524
init 2 7 13240
kill 2 8 12708
ln 2 9 12604
ls 2 10 14792
mkdir 2 11 12788
rm 2 12 12764
sh 2 13 23252
stressfs 2 14 13436
usertests 2 15 56364
wc 2 16 14184
zombie 2 17 12428
mytest 2 18 12612
console 3 19 0
$ mytest
1: init
2: sh
3: Wrong pid
4: mytest
5: Wrong pid
6: Wrong pid
7: Wrong pid
8: Wrong pid
9: Wrong pid
10: Wrong pid
$ 
```

Submission

- This project is to implement only the system calls (getnice, setnice, ps)
 - The user program for testing is irrelevant.
- Use the ***submit*** & ***check-submission*** binary file in Ui Server
 - **\$ make clean**
 - **\$ ~swe3004/bin/submit pa1 xv6-public**
 - You can submit several times, and the submission history can be checked through check-submission
 - Only the last submission will be graded

Submission

- PLEASE DO NOT COPY
 - We will run inspection program on all the submissions
 - Any unannounced penalty can be given to **both students**
 - 0 points / negative points / F grade ...
- Due date: 10/2(Wed.), 23:59:59 PM
 - -25% per day for delayed submission

Questions

- If you have questions, please ask on i-campus
 - Please use the discussion board
 - Discussion board preferred over messages
- You can also visit Corporate Collaboration Center #85533
 - Please iCampus message TA before visiting
- Reading xv6 commentary will help you a lot
 - <http://csl.skku.edu/uploads/SSE3044S20/book-rev11.pdf>

Appendix. Trap Handling Process on xv6

- Example : kill system call

```
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 int
6 main(int argc, char **argv)
7 {
8     int i;
9
10    if(argc < 2){
11        printf(2, "usage: kill pid...\n");
12        exit();
13    }
14    for(i=1; i<argc; i++)
15        kill atoi(argv[i]);
16    exit();
17 }
```

kill.c (user level)

```
1 struct stat;
2 struct rtcdate;
3
4 // system calls
5 int fork(void);
6 int exit(void) __attribute__((noreturn));
7 int wait(void);
8 int pipe(int*);
9 int write(int, const void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int); // Kill system call
13 int exec(char*, char**);
14 int open(const char*, int);
15 int mknod(const char*, short, short);
16 int unlink(const char*).
```

user.h

_kill(user program)'s Build Process

```
ld -m elf_i386 -N -e main -Ttext 0 -o _cat cat.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _echo echo.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _forktest forktest.o ulib.o usys.o
ld -m elf_i386 -N -e main -Ttext 0 -o _grep grep.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _init init.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _kill kill.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _ln ln.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _ls ls.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _mkdir mkdir.o ulib.o usys.o printf.o umalloc.o
ld -m elf_i386 -N -e main -Ttext 0 -o _rm rm.o ulib.o usys.o printf.o umalloc.o
```

make qemu-nox | grep usys

Functions defined as assembly

```
1 #include "syscall.h"
2 #include "traps.h"
3
4 #define SYSCALL(name) \
5     .globl name; \
6     name: \
7         movl $SYS_ ## name, %eax; \
8         int $T_SYSCALL; \
9         ret
10
11 SYSCALL(fork)
12 SYSCALL(exit)
13 SYSCALL(wait)
14 SYSCALL(pipe)
15 SYSCALL(read)
16 SYSCALL(write)
17 SYSCALL(close)
18 SYSCALL(kill)
19 SYSCALL(exec)
```

usys.S

```
1 #include "syscall.h"
2
3 // These are arbitrarily chosen
4 // processor defined exceptions
5 #define SYS_pipe        4
6 #define SYS_read         5
7 #define SYS_kill         6
8 #define SYS_exec         7
9 #define SYS_fstat        8
```

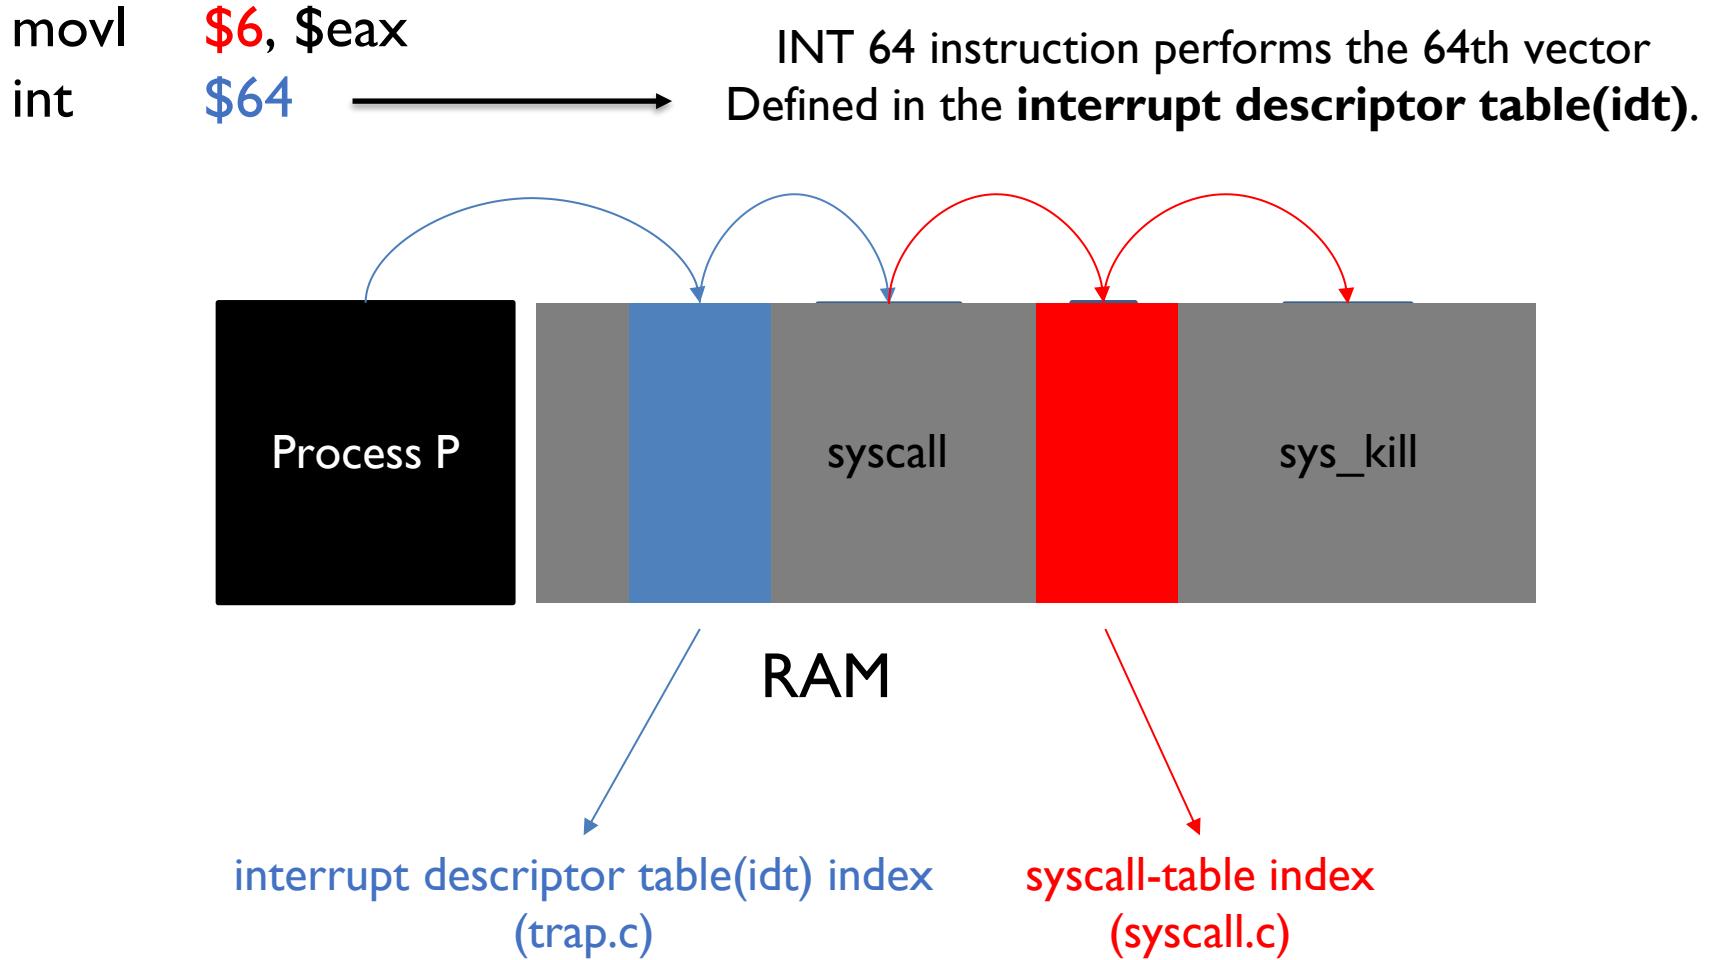
syscall.h

```
25 // These are arbitrarily chosen
26 // processor defined exceptions
27 #define T_SYSCALL        64
28 #define T_DEFAULT         500
```

traps.h

After all, what the `_kill` user program calls is three instructions.

Trap Handling Process on xv6



Interrupt Descriptor Table (IDT)

main.c

```
// Bootstrap processor starts running C code here.
// Allocate a real stack and switch to it, first
// doing some setup required for memory allocator to work.
int
main(void)
{
    kinit1(end, P2V(4*1024*1024)); // phys page allocator
    kvmalloc(); // kernel page table
    mpinit(); // detect other processors
    lapicinit(); // interrupt controller
    seginit(); // segment descriptors
    picinit(); // disable pic
    ioapicinit(); // another interrupt controller
    consoleinit(); // console hardware
    uartinit(); // serial port
    pinit(); // process table
    tvinit(); // trap vectors
    buncle(); // buffer cache
    fileinit(); // file table
    ideinit(); // disk
    startothers(); // start other processors
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
    userinit(); // first user process
    mpmain(); // finish this processor's setup
}
```

https://wiki.osdev.org/Interrupt_Descriptor_Table

IDT entry, Interrupt Gates				Description																				
Name	Bit	Full Name	Offset																					
Offset	48..63	Offset 16..31		Higher part of the offset.																				
P	47	Present		Set to 0 for unused interrupts.																				
DPL	45..46	Descriptor Privilege Level		Gate call protection. Specifies which privilege Level the calling Descriptor minimum should have. So hardware and CPU interrupts can be protected from being called out of userspace.																				
S	44	Storage Segment		Set to 0 for interrupt and trap gates (see below).																				
Type	40..43	Gate Type 0..3		Possible IDT gate types :																				
				<table border="1"> <tr><td>0b0101</td><td>0x5</td><td>5</td><td>80386 32 bit task gate</td></tr> <tr><td>0b0110</td><td>0x6</td><td>6</td><td>80286 16-bit interrupt gate</td></tr> <tr><td>0b0111</td><td>0x7</td><td>7</td><td>80286 16-bit trap gate</td></tr> <tr><td>0b1110</td><td>0xE</td><td>14</td><td>80386 32-bit interrupt gate</td></tr> <tr><td>0b1111</td><td>0xF</td><td>15</td><td>80386 32-bit trap gate</td></tr> </table>	0b0101	0x5	5	80386 32 bit task gate	0b0110	0x6	6	80286 16-bit interrupt gate	0b0111	0x7	7	80286 16-bit trap gate	0b1110	0xE	14	80386 32-bit interrupt gate	0b1111	0xF	15	80386 32-bit trap gate
0b0101	0x5	5	80386 32 bit task gate																					
0b0110	0x6	6	80286 16-bit interrupt gate																					
0b0111	0x7	7	80286 16-bit trap gate																					
0b1110	0xE	14	80386 32-bit interrupt gate																					
0b1111	0xF	15	80386 32-bit trap gate																					
0	32..39	Unused 0..7		Have to be 0.																				
Selector	16..31	Selector 0..15		Selector of the interrupt function (to make sense - the kernel's selector). The selector's descriptor's DPL field has to be 0 so the iret instruction won't throw a #GP exception when executed.																				
Offset	0..15	Offset 0..15		Lower part of the interrupt function's offset address (also known as pointer).																				

In xv6, idt is set in the form shown in the Intel architecture manual.

trap.c

```
17 void
18 tvinit(void)
19 {
20     int i;
21
22     for(i = 0; i < 256; i++)
23         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
24     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL] DPL_USER);
25
26     initlock(&tickslock, "time");
27 }
```

mmu.h

```
// Gate descriptors for interrupts and traps
struct gatedesc {
    uint off_15_0 : 16; // low 16 bits of offset in segment
    uint cs : 16; // code segment selector
    uint args : 5; // # args, 0 for interrupt/trap gates
    uint rsv1 : 3; // reserved(should be zero I guess)
    uint type : 4; // type(STS_IG32,TG32)
    uint s : 1; // must be 0 (system)
    uint dpl : 2; // descriptor(meaning new) privilege level
    uint p : 1; // Present
    uint off_31_16 : 16; // high bits of offset in segment
};

// Set up a normal interrupt/trap gate descriptor.
// - istrap: 1 for a trap (= exception) gate, 0 for an interrupt gate.
// - interrupt gate clears FL_IF, trap gate leaves FL_IF alone
// - sel: Code segment selector for interrupt/trap handler
// - off: Offset in code segment for interrupt/trap handler
// - dpl: Descriptor Privilege Level -
//       the privilege level required for software to invoke
//       this interrupt/trap gate explicitly using an int instruction.
#define SETGATE(gate, istrap, sel, off, d)
{
    (gate).off_15_0 = (uint)(off) & 0xffff;
    (gate).cs = (sel);
    (gate).args = 0;
    (gate).rsv1 = 0;
    (gate).type = (istrap) ? STS_TG32 : STS_IG32;
    (gate).s = 0;
    (gate).dpl = (d);
    (gate).p = 1;
    (gate).off_31_16 = (uint)(off) >> 16;
}
```

Vector

As a result, int 64 calls vector64, and vector64 executes alltraps.

```
316    jmp alltraps  
317 .globl vector64  
318 vector64:  
319     pushl $0  
320     pushl $64  
321     jmp alltraps  
322 .globl vector65  
323 vector65:
```

vectors.S

```
5 alltraps:  
6     # Build trap frame.  
7     pushl %ds  
8     pushl %es  
9     pushl %fs  
10    pushl %gs  
11    pushal  
12  
13    # Set up data segments.  
14    movw $(SEG_KDATA<<3), %ax  
15    movw %ax, %ds  
16    movw %ax, %es  
17  
18    # Call trap(tf), where tf=%esp  
19    pushl %esp  
20    call trap  
21    addl $4, %esp  
22  
23    # Return falls through to trapret
```

trapasm.S

Trap

```
36 void
37 trap(struct trapframe *tf)
38 {
39     if(tf->trapno == T_SYSCALL){
40         if(myproc()->killed)
41             exit();
42         myproc()->tf = tf;
43         syscall();
44         if(myproc()->killed)
45             exit();
46     }
47 }
```

trap.c

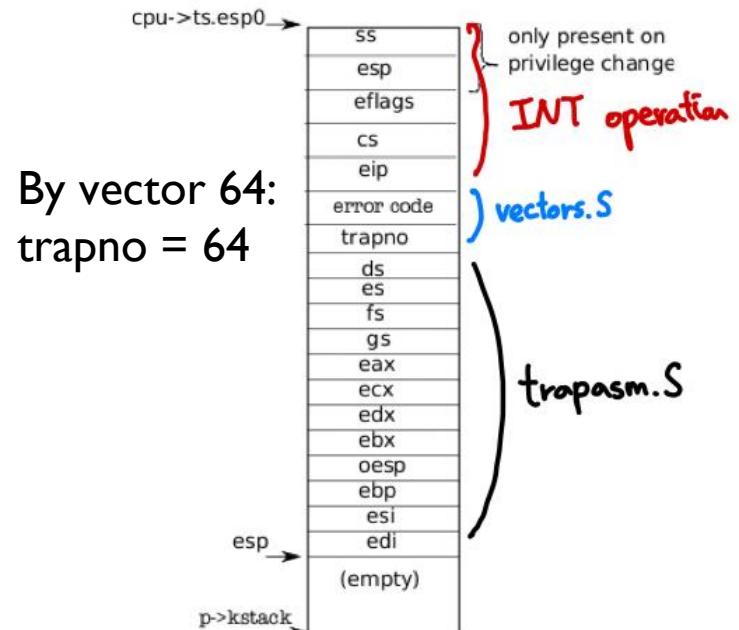


Figure 2-2. The trapframe on the kernel stack

Syscall

```
139 void
140 syscall(void)
141 {
142     int num;
143     struct proc *curproc = myproc();
144
145     num = curproc->tf->eax;          movl    $6, $eax
146     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {    int      $64
147         curproc->tf->eax = syscalls[num]();                      ; Call to syscalls[num]
148     } else {
149         cprintf("%d %s: unknown sys call %d\n",
150                 curproc->pid, curproc->name, num);
151         curproc->tf->eax = -1;
152     }
153 }
```

syscall.c

```
111 static int (*syscalls[])(void) = {
112     [SYS_fork]    sys_fork,
113     [SYS_exit]   sys_exit,
114     [SYS_wait]    sys_wait,
115     [SYS_pipe]   sys_pipe,
116     [SYS_read]   sys_read,
117     [SYS_kill]    sys_kill,    ; Corresponds to num=6 in the code
118     [SYS_exec]   sys_exec,
119     [SYS_fstat]  sys_fstat.
```

Kill

```
29 int
30 sys_kill(void)
31 {
32     int pid;
33
34     if(argint(0, &pid) < 0)
35         return -1;
36     return kill(pid);
37 }
```

sysproc.c

```
479 int
480 kill(int pid)
481 {
482     struct proc *p;
483
484     acquire(&ptable.lock);
485     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
486         if(p->pid == pid){
487             p->killed = 1;
488             // Wake process from sleep if necessary.
489             if(p->state == SLEEPING)
490                 p->state = RUNNABLE;
491             release(&ptable.lock);
492             return 0;
493         }
494     }
495     release(&ptable.lock);
496     return -1;
497 }
```

proc.c